

■ SFT Interview Prep – Infographic Style

■ What is SFT?

- ➔■ Supervised Fine-Tuning = train pre-trained LLM on **labeled input–output pairs**
- ➔■ First step before DPO / RLHF
- ➔■ Goal: **align model behavior** (instruction following, domain knowledge, safety)

■ Top Use Cases

1■ Instruction Following	■ JSON formatting, summarization, style control
2■ Domain Adaptation	■ Medical, Legal, Finance, E-commerce
3■ Samsung E-commerce Example	■ Example understanding → “budget phone with good camera” → Galaxy A series ■ Product recommendation
4■ Policy & Safety	■ Avoid disallowed content, enforce rules

■ How SFT Changes Model Behavior

Before SFT	After SFT
■ Generic responses ■ Not domain-specific	■ Domain-specific (Samsung products) ■ Consistent format (JSON)

■ Interview Quick Points

- Fastest + cheapest way to specialize a model
- Needs **high-quality labeled data** (search logs, FAQs, QA pairs)
- Risk: too narrow = catastrophic forgetting
- SFT sets **foundation**, DPO/RLHF refine preferences

■ Elevator Answer

“SFT fine-tunes LLMs on labeled pairs to align them with tasks. For Samsung e-commerce, it helps rank relevant products, answer FAQs, and map vague queries like *‘best phone for gaming’* to the correct SKUs. It’s the fastest way to make a base LM domain-ready.”