

HYBRID SEARCH

Optimizing the R in RAG

Apr 16, 2025

Can't cover in 45 mins...

1. How lexical search actually works (ask chat GPT about: inverted index, read “Relevant Search” 😊)
2. What is an embedding
3. Lexical scoring, vector scoring (cosine, euclidean, etc similarities) etc


Intuitive sense of “close” good enough for today :)

Traditional Keyword Search



VS

Vector Search



Method

Matches exact words or phrases

Transforms text, audio, and images into numeric representations and matches based on meaning, intent, and context

Speed

May slow down with large and varied data

Fast, even on large and varied data

Relevance of Results

Limited by exact word or phrase match

Highly relevant due to understanding of context and meaning

Flexibility

Primarily text-based search

Enables new classes of search for text, image, and audio

	Traditional (Keyword) Search	Vector Search
SEARCH APPROACH	Keyword-based matching	Semantic/meaning-based matching using vector embeddings
AMBIGUITY HANDLING	Struggles with synonyms, ambiguous language, fuzzy queries	Handles synonyms, ambiguous language, fuzzy queries
SEARCH RELEVANCE CALCULATION	BM25, TF-IDF	Jaccard Similarity, Cosine Similarity, and L2 Distance
SEARCH QUALITY	Depends on exact keyword matching	Semantic relationships => better handles broad queries
SPEED AND IMPLEMENTATION	Fast for simple queries, easy implementation, straightforward usage	Slightly slower, more complicated to implement (if ANN is used)
SCALIBILITY	Challenged by the continuous expansion of content	Better handles huge datasets
COST	Low computational requirements => lower cost	High computational requirements (without ANN) => higher cost

Traditional search vs. vector search

	Traditional (Keyword) Search	Vector Search
SEARCH APPROACH	Matches the exact keywords.	Finds related objects that share similar characteristics.
AMBIGUITY HANDLING	Struggles with ambiguous language and synonyms.	Uses machine learning models to handle synonyms and ambiguous language.
RELEVANCE	Works well for precise queries.	Superior for broad or fuzzy queries.
SPEED AND SCALABILITY	Fast, scales well due to simple index reading.	Less efficient and struggles with scaling due to complex vector calculations.
SEARCH QUALITY	Mostly depends on exact keyword matching.	Uses semantic relationships, thus better handling of broad queries.
COST	Generally less due to lower computational requirement.	Usually higher due to the need for more computing power.

Lexical Search

Key Points:

- **Definition:** Finds documents with exact query words.
- **Core Structure:** Inverted Index \rightarrow term \rightarrow list of doc IDs.
 - Example: "apple" \rightarrow [Doc2, Doc5, Doc9]
- **Process:**
 - Tokenize query
 - Lookup tokens in index
 - Retrieve matching docs
 - Apply scoring
- **Scoring Methods:**
 - **TF-IDF:** Weights term frequency \times rarity
 - **BM25:** Length normalization + term frequency saturation
- **Pros:** Fast, explainable
- **Cons:** No semantic understanding ("phone" \neq "mobile")

Vector Search

Key Points:

- **Definition:** Matches based on semantic meaning, not exact words.
- **Embedding:** Vector representation of meaning
 - Similar meaning → vectors close together
 - Example: "dog" \approx "puppy", far from "laptop"
- **Embedding Generation:** Word2Vec, GloVe, Transformers, etc.
- **Vector Scoring Methods:**
 - **Cosine Similarity:** Angle between vectors
 - **Euclidean Distance:** Straight-line distance
 - **Dot Product:** Magnitude + direction
- **Pros:** Finds synonyms, concept matches ("TV" \approx "television")
- **Cons:** Needs embedding computation, slower than lexical search

Also won't cover

1. RRF – Reciprocal Rank Fusion

RRF is Not Enough

NOVEMBER 3RD, 2024

Hybrid search means combining lexical and vector search results into one result listing.

“We’ll just use [Reciprocal Rank Fusion](#)” I’m sure I’ve said from time to time.

As if RRF is kind of “a miracle occurs”. You get the best of both worlds, and suddenly your search looks incredible.

Take the query `hello to the planet`. Let’s say we start with reasonable results from a vector search system (follow along in [this notebook](#))

<code>vector_sim</code>	<code>texts</code>	<code>vector_rank</code>
-------------------------	--------------------	--------------------------

Assumption: embeddings good first pass search

Embeddings get you *close* but not all the way


ID	Title	Vector (256? 512? Or more dimensions)
0	mary had a little lamb	[0.9, 0.8, -0.5, 0.75, ..]
1	mary had a little ham	[0.6, 0.4, -0.4, 0.60, ..]
2	a little ham	[-0.2, 0.5, 0.9, -0.45, ..]
3	little mary had a scam	[0.4, -0.5, 0.25, 0.14, ..]
4	ham it up with mary	[0.2, 0.5, 0.2, 0.45, ..]
5	Little red riding hood had a baby sheep?	[0.95, 0.79, -0.49, 0.65, ..]

Similar!

(despite sharing few terms)

Chunked

You've chunked your data into a meaningful "search document" with important metadata:



```
{  
  "Book_title": "Nursery Rhymes"  
  "Section": "Mary Had a Little  
  Lamb"  
  "Text": "..."  
}
```

Embedding for *whole document*

We want an embedding capturing as much of the document as is reasonable

```
text_concatted = data['product_name'] + ' -- ' + data['product_description']  
embedding = model.encode(text_concatted)
```

(Not just a title embedding)

Embedding is ~ two-towerable

Short text (ie queries) and long text (paragraphs) can be mapped in similarity space

QUERY: Kid story
about sheep

Document:

Mary had a little lamb, little
lamb, little lamb.

Mary had a little lamb, its
fleece was white as snow.

And everywhere that Mary went.
Mary went. Mary went.

And everywhere that Mary went,
the lamb was sure to go.

It followed her to school one
day, school one day, school one




Similar

The diagram illustrates the concept of mapping text into a similarity space. On the left, a query 'Kid story about sheep' has an arrow pointing towards the center. On the right, a document containing several paragraphs of the 'Mary Had a Little Lamb' story has an arrow pointing towards the center. In the center, the word 'Similar' is displayed, indicating the result of the mapping process.

Bonus: embedding is a two tower model!


Query Features

- Query embedding
- Query



Document Features

- Name
- Description
- Product image embedding
- ???



(Biencoder,
learned on
labeled data)

After embedding we boost/rerank/...

Exact name match?

- Move these to the top!

Query mentions color?

- Ensure color matches boosted

Query Understanding

Home About


☆ Pinned

 Daniel Tunkelang

Query Understanding: An Introduction

Search engines are so core to our digital experience that we take them for granted. Most of us cannot remember the web without...


Dec 2, 2023 🗨️ 242



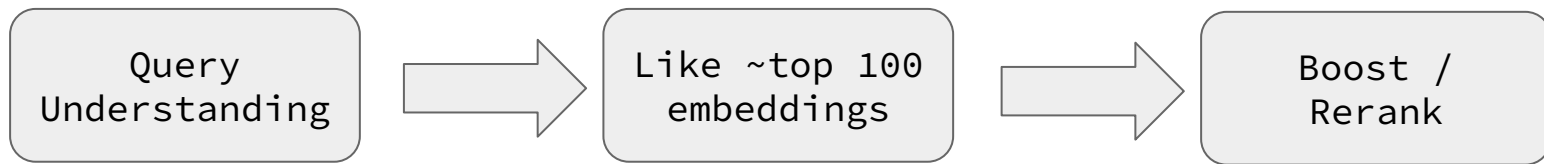
(Different query types == different treatments!)

<http://queryunderstanding.com>

Ideal:



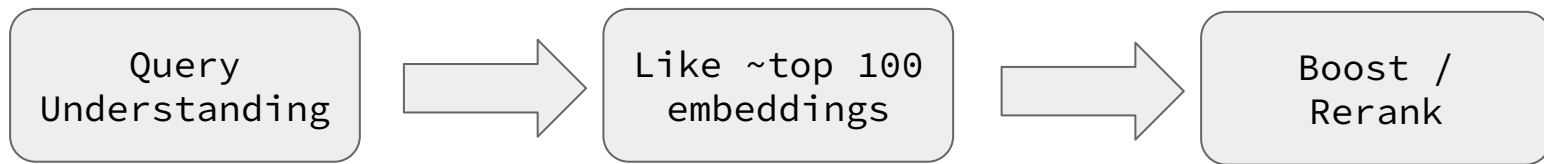
Reality:



Reality:



Reality:



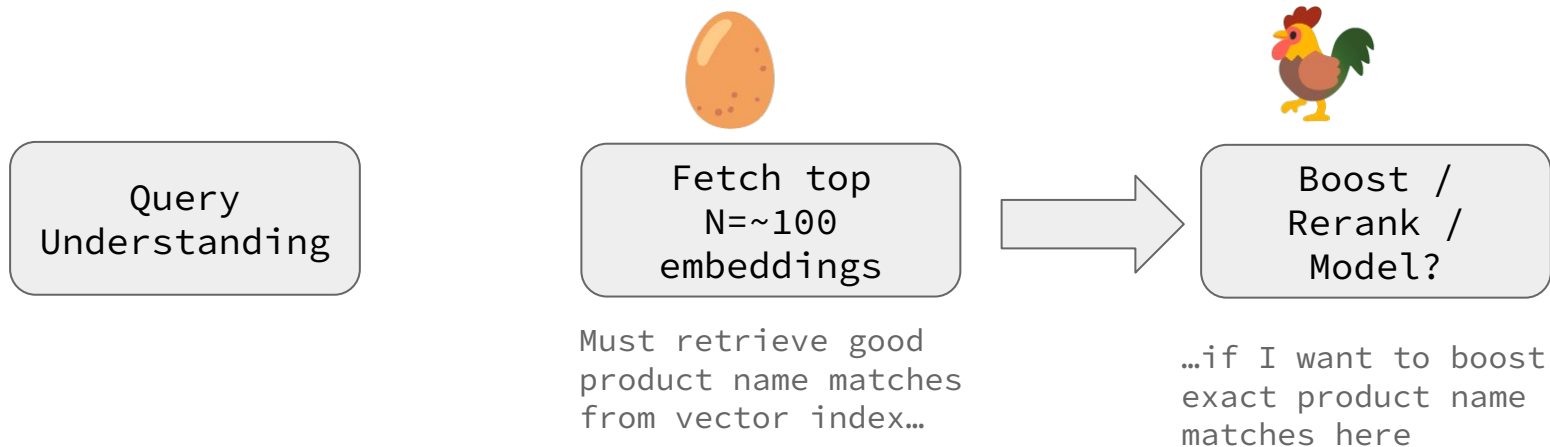
Need to filter this to
“good” 100 or so

Chicken and egg problem:

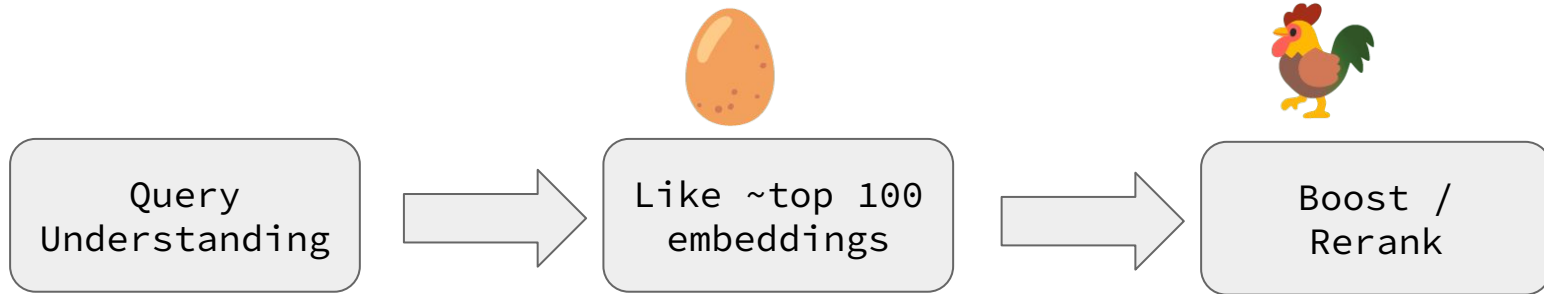


If I want to boost
exact product name
matches here..

Chicken and egg problem:



Chicken and egg problem:



The good product name
matches better be in
the candidates!

~2021 vector DB

No WHERE!

```
SELECT * FROM <search_engine>
```



Can't guarantee
product name matches
promoted

```
ORDER BY vector_similarity(query_embedding, title_embedding)  
LIMIT 100
```

2025 vector DB (search engine)

```
SELECT * FROM <search>
```

```
WHERE [trowel] in product_name
```

BEFORE vector_similarity
Get candidates matching
“trowel”



Now I have matches!

...

```
ORDER BY vector_similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

~2025 era vector DB (search engine)

```
SELECT * FROM <search>
```

```
WHERE [trowel] in product_name
```

```
...
```

```
ORDER BY vector_similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

BEFORE vector_similarity
Get candidates matching
“mary”



*How does your vector DB
pre-filter? Can you do this
at scale?*

... and “where” could be *anything*

Search for “garden trowel”

```
SELECT * FROM <search>
```

```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```

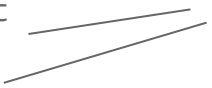
```
AND (garden in title OR garden in description OR
```

```
    trowel in title OR trowel in description)
```

```
ORDER BY vector_similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

Somehow we turn the query
to this dept / item type



... and “where” could be *anything*

Search for “garden trowel”

```
SELECT * FROM <search>
```

```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```


```
AND (garden in title OR garden in description OR
```

```
trowel in title OR trowel in description)
```

```
ORDER BY vector_similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

And also match
query terms in
tokenized
title/description



... and “where” could be *anything*

Search for “garden trowel”

```
SELECT * FROM <search>
```

```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```

```
AND (garden in title OR garden in description OR
```

```
    trowel in title OR trowel in description)
```

```
ORDER BY vector_similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

And also match
query terms

*(yes you search nerds,
I’m ignoring BM25 and
lexical scoring for now)*

Practically: there's a vector index

We can reasonably get top K...

Search for “garden trowel”

```
SELECT * FROM <search>
```

```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```

```
AND (garden in title OR garden in description OR
```

```
trowel in title OR trowel in description)
```

```
ORDER BY vector_similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

Get top 100 from
this set via an
index

(otherwise we scan all
results to score them)

There's more than one “top K” we care about

What about “pure” vector matches?

```
SELECT * FROM <search>
```

```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```

```
AND (garden in title OR garden in description OR
```

```
trowel in title OR trowel in description)
```

```
ORDER BY similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

100 from this set

UNION ALL

```
SELECT * FROM <search>
```

```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```

```
ORDER BY similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

There's more than one candidate set

What about “pure” vector matches?

```
SELECT * FROM <search>
```

```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```

```
AND (garden in title OR garden in description OR
```

```
trowel in title OR trowel in description)
```

```
ORDER BY similarity(query_embedding, title_embedding)
```

```
LIMIT 100
```

UNION ALL

+ 100 from this set

```
SELECT * FROM <search>
```


```
WHERE “lawn_and_garden” in department
```

```
AND “trowel” in item_type
```

```
ORDER BY similarity(query_embedding, title_embedding)
```


```
LIMIT 100
```

With squiggly lines...

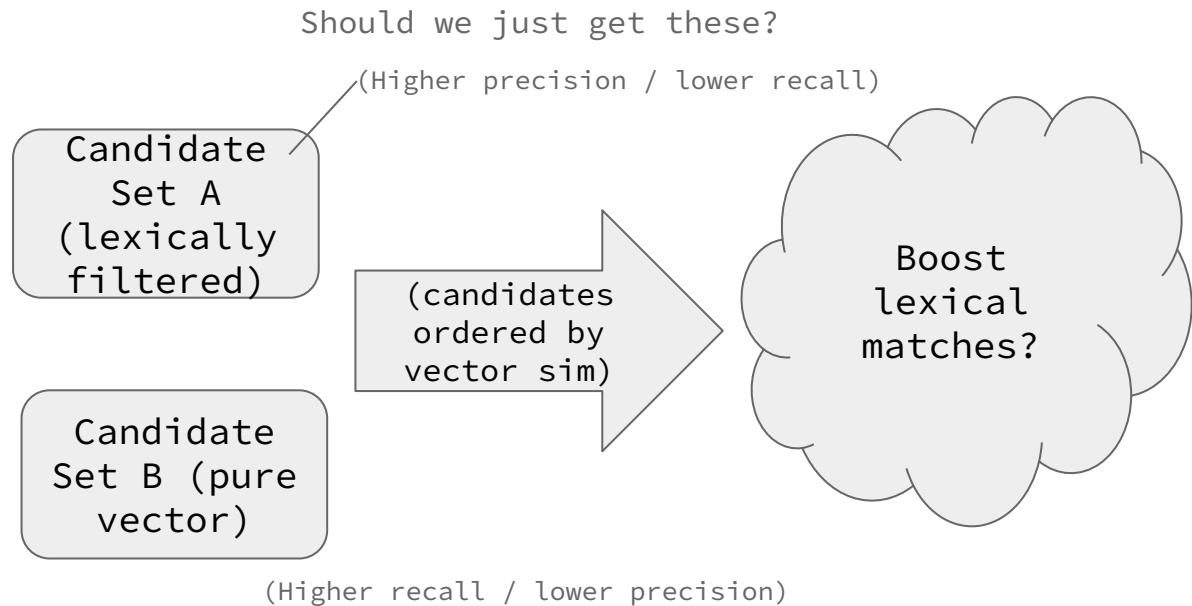


Why do we do it this way?

Should we just get these?



Why do we do it this way?



With squiggly lines...

L0 Retrieval

Candidate
Set A
(filtered
to lexical)

Candidate
Set B (pure
vector)

(candidates
ordered by
vector sim)

L1 Ranking

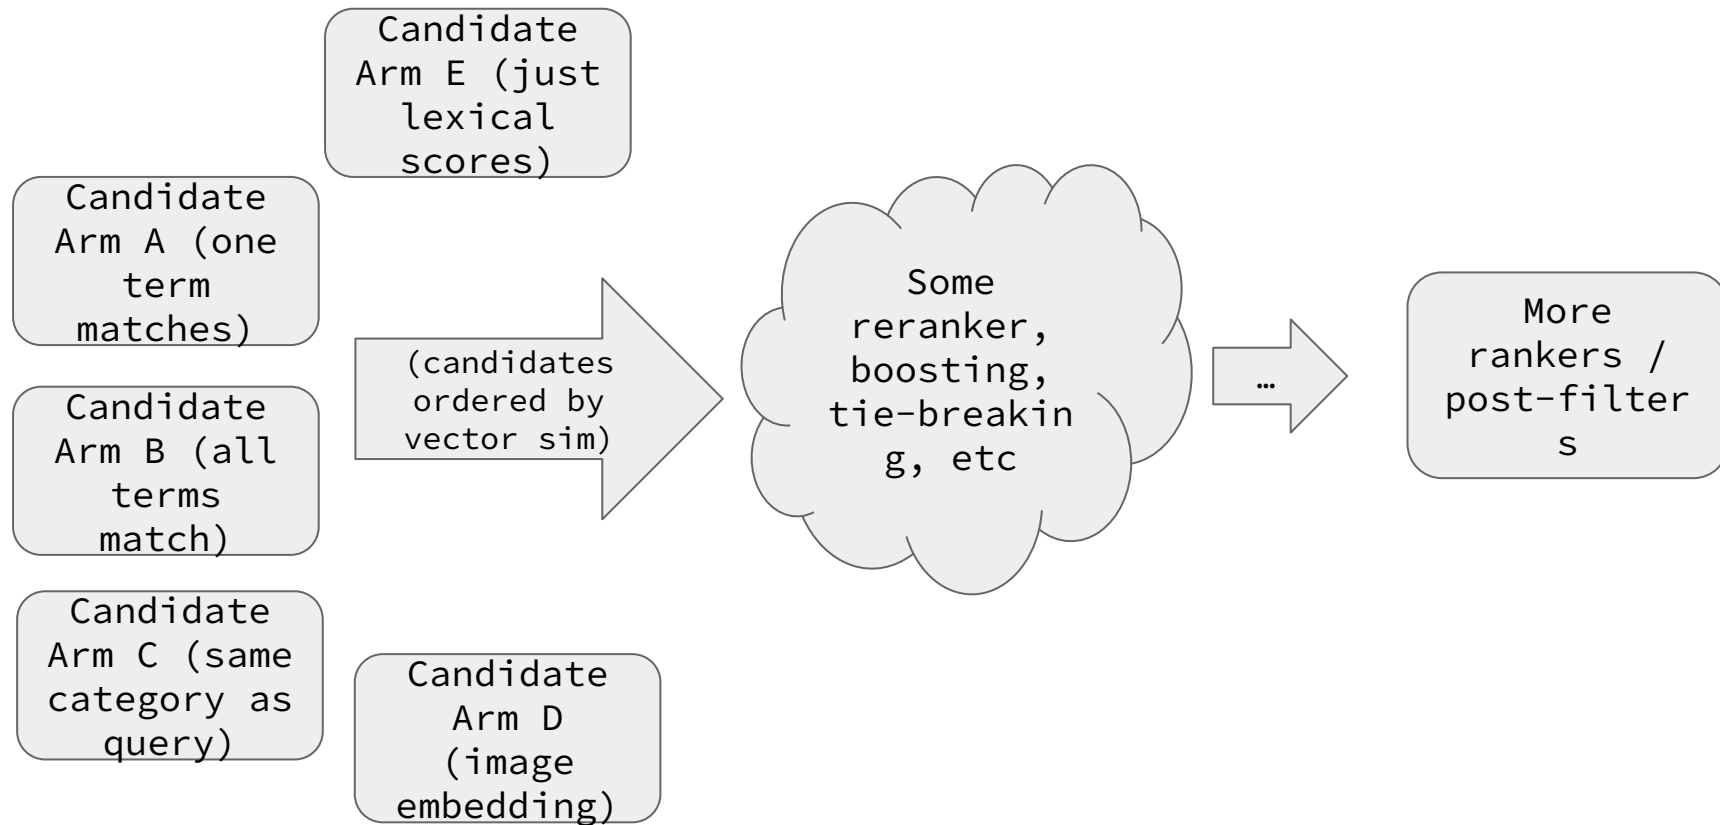
Some
reranker,
boosting,
tie-breakin
g, etc

...

More
rankers /
post-filter
s

A retrieval “Arm”

And many retrieval arms



L0

Retrieval Arms



Candidate
Arm E (just
lexical
scores)

Candidate
Arm A (one
term
matches)

Candidate
Arm B (all
terms
match)

Candidate
Arm C (same
category as
query)

Candidate
Arm D
(image
embedding)

(candidates
ordered by
vector sim)

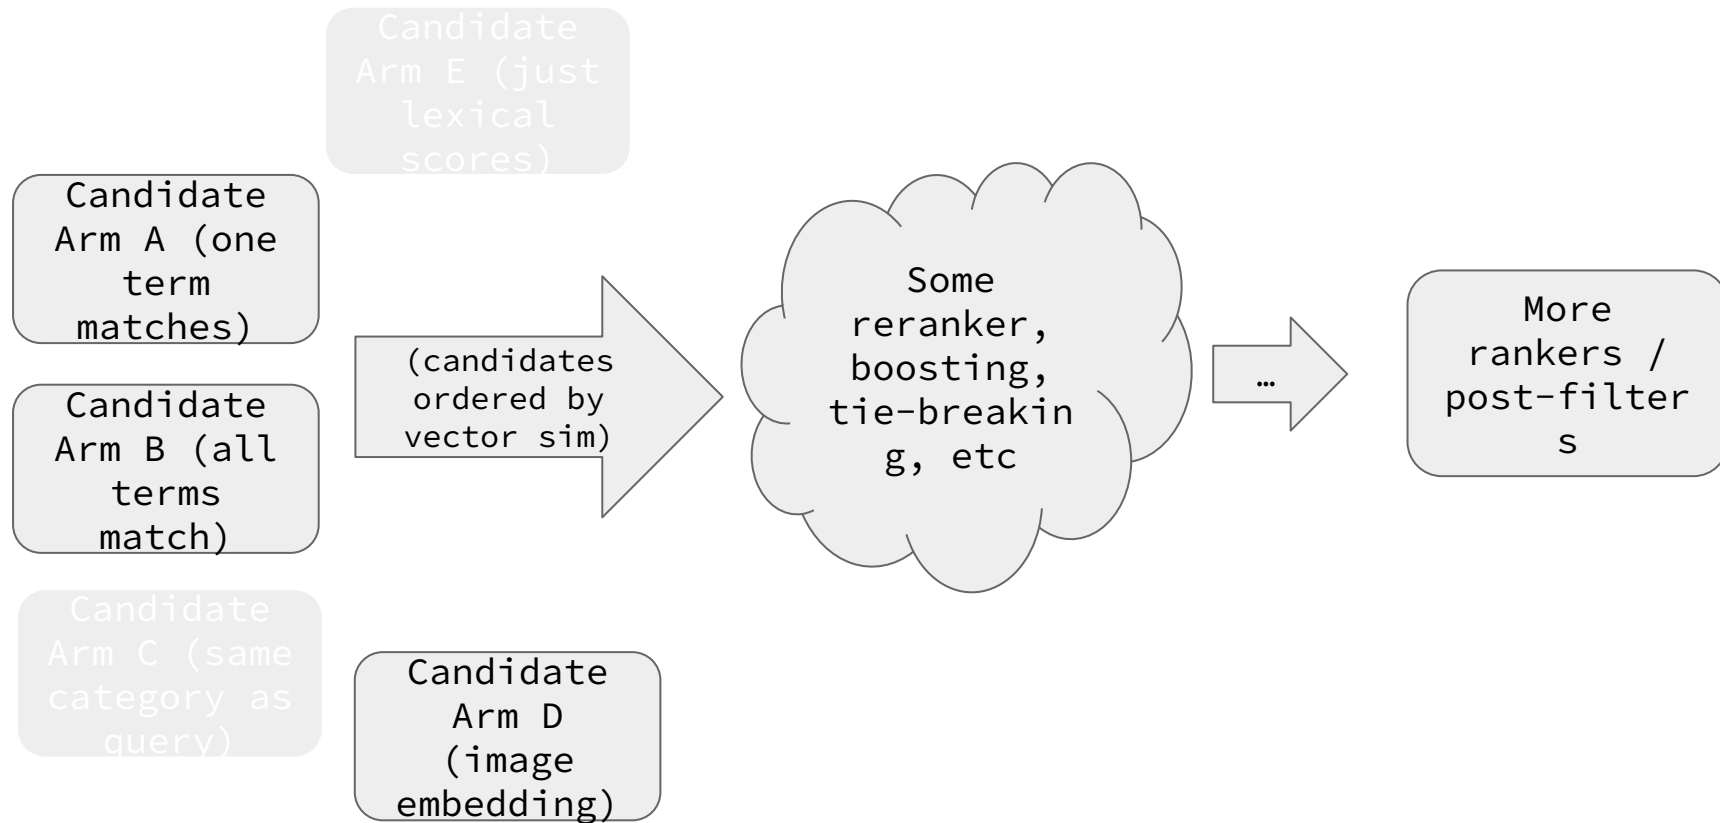
L1

boost/reranking




Boost /
Rerank

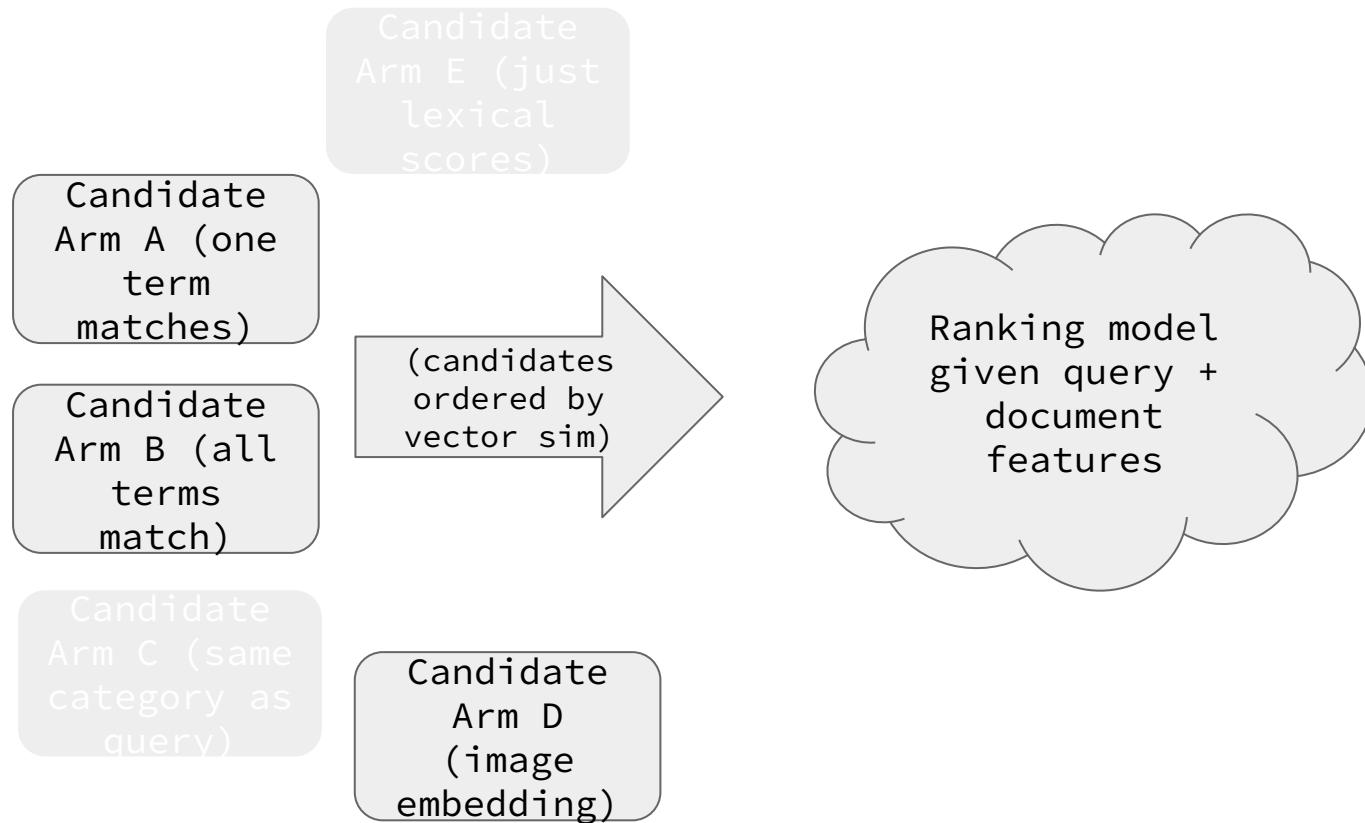
Or depending on the query



Then the boost



Or a model



That's the theory at least

